

Language Representation

Lecturer: Eilam Shapira



NLP Course - Unit 3 - Lecture

Based on Amir Feder's slides & Stanford's CS224n

Today's Agenda

- **Word2Vec**
- GloVe
- Sentence representations
- Applications

Before we start...

- Can you engineer features to represent:
 - Words?
 - Sentences?
 - Documents?
 - What about yourselves?
- What's the dimension?
 - Remember PCA?

Embedding is a powerful way to represent data

Openness to experience 79 out of 100

Agreeableness 75 out of 100

Conscientiousness 42 out of 100

Negative emotionality 50 out of 100

Extraversion 58 out of 100



10. Lionel Messi
Attacking Mid. (Right), Attacking Mid. (Center), Striker (Center)

 **32 years old (24/6/1987)**
Main 136 caps / 68 goals
Youth 18 caps / 14 goals
170 cm / 5' 6" | 72 kg

 Contracted to **FC Barcelona**
Valued at € 85.2M
€ 71.3M p/a until 30/6/2021
Legendary attacking midfielder

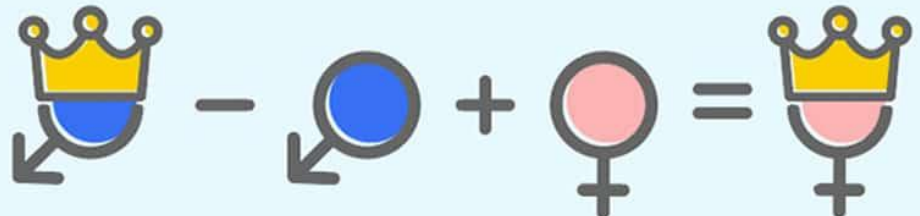
| TECHNICAL | MENTAL | PHYSICAL |
|---------------------|------------------|---|
| Corners 15 | Aggression 7 | Acceleration 18 |
| Crossing 15 | Anticipation 19 | Agility 19 |
| Dribbling 20 | Bravery 10 | Balance 19 |
| Finishing 20 | Composure 18 | Jumping Reach 6 |
| First Touch 19 | Concentration 13 | Natural Fitness 14 |
| Free Kick Tak... 19 | Decisions 20 | Pace 15 |
| Heading 10 | Determination 20 | Stamina 13 |
| Long Shots 17 | Flair 20 | Strength 9 |
| Long Throws 4 | Leadership 14 | Foot Left |
| Marking 4 | Off The Ball 16 |  |
| Passing 20 | Positioning 5 | |
| Penalty Taking 17 | Teamwork 14 | |
| Tackling 7 | Vision 20 | |
| Technique 20 | Work Rate 7 | |

FM2020 | (March 6, 2020)

fmdataba.com
football manager database

Similarity-preserving Representation

Word2vec



The diagram shows a vector equation: King - Man + Woman = Queen. Each word is represented by an icon. 'King' is a yellow crown on a blue circle with a male symbol. 'Man' is a blue circle with a male symbol. 'Woman' is a pink circle with a female symbol. 'Queen' is a yellow crown on a pink circle with a female symbol. The equation is represented by minus, plus, and equals signs.

King Man Woman Queen

Word2Vec

Word2vec: Skip-Gram Task

- Word2vec provides a variety of options. Let's do
 - "skip-gram with negative sampling" (SGNS)
 - ~~○ Continuous Bag-of-words (CBOW)~~
- Word2vec actually uses a logistic regression initialized with random weights
 - But we'll show it with a shallow NN

Skip-Gram Algorithm

1. Treat the target word and a neighboring context word as positive examples.
2. Randomly sample other words in the lexicon to get negative samples
3. Use NN to train a classifier to distinguish those two cases
4. Use the weights as the embeddings

Skip-Gram Goal

Given a tuple (t,c) = target, context

- (*hotel*, *haifa*)
- (*hotel*, *aardvark*)

Return probability that c is a real context word:

$$P(+ | t, c)$$

$$P(- | t, c) = 1 - P(+ | t, c)$$

How to compute $p(+ | t, c)$?

Intuition:

- Words are likely to appear near similar words
- Model similarity with dot-product!
- $\text{Similarity}(t, c) \propto t \cdot c$

Problem:

- *Dot product is not a probability!*
 - *(Neither is cosine)*

Turning dot product into a probability

$$P(+|t, c) = \frac{1}{1+e^{-t \cdot c}}$$

$$\begin{aligned} P(-|t, c) &= 1 - P(+|t, c) \\ &= \frac{e^{-t \cdot c}}{1+e^{-t \cdot c}} \end{aligned}$$

For all the context words

Assume they are independent

$$P(+|t, c_{1:k}) = \prod_{i=1}^K \frac{1}{1+e^{-t \cdot c_i}}$$

$$\log P(+|t, c_{1:k}) = \sum_{i=1}^K \log \frac{1}{1+e^{-t \cdot c_i}}$$

Positive Training Examples

Source Text

Training Samples

The quick brown fox jumps over the lazy dog. ➡

(the, quick)
(the, brown)

The quick brown fox jumps over the lazy dog. ➡

(quick, the)
(quick, brown)
(quick, fox)

The quick brown fox jumps over the lazy dog. ➡

(brown, the)
(brown, quick)
(brown, fox)
(brown, jumps)

The quick brown fox jumps over the lazy dog. ➡

(fox, quick)
(fox, brown)
(fox, jumps)
(fox, over)

Negative Training Examples

| Source Text | Training Samples | Negative Samples | | | | | | |
|---|------------------|------------------|-------|------------------------------|--|--|---|------|
| <table><tr><td>The</td><td>quick</td><td>brown</td></tr></table> fox jumps over the lazy dog. ➡ | The | quick | brown | (the, quick) (the, brown) | (the, aardvark) ... | | | |
| The | quick | brown | | | | | | |
| <table><tr><td>The</td><td>quick</td><td>brown</td><td>fox</td></tr></table> jumps over the lazy dog. ➡ | The | quick | brown | fox | (quick, the) (quick, brown) (quick, fox) | (the,zztop) | | |
| The | quick | brown | fox | | | | | |
| <table><tr><td>The</td><td>quick</td><td>brown</td><td>fox</td><td>jumps</td></tr></table> over the lazy dog. ➡ | The | quick | brown | fox | jumps | (brown, the) (brown, quick) (brown, fox) (brown, jumps) | (brown, aardvark) ... (brown,zztop) | |
| The | quick | brown | fox | jumps | | | | |
| <table><tr><td>The</td><td>quick</td><td>brown</td><td>fox</td><td>jumps</td><td>over</td></tr></table> the lazy dog. ➡ | The | quick | brown | fox | jumps | over | (fox, quick) (fox, brown) (fox, jumps) (fox, over) | |
| The | quick | brown | fox | jumps | over | | | |

Objective Criteria

We want to maximize...

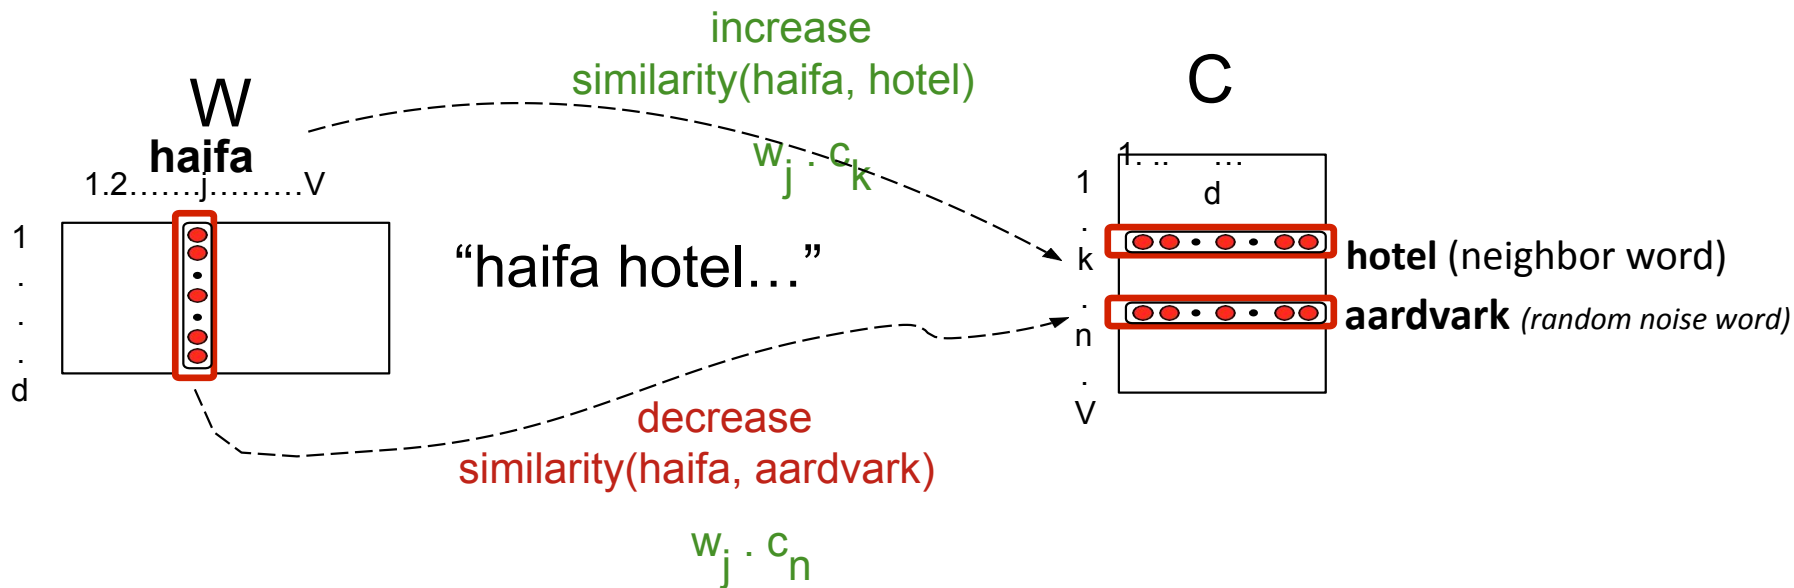
$$\sum_{(t,c) \in +} \log P(+|t, c) + \sum_{(t,c) \in -} \log P(-|t, c)$$

Maximize the + label for the pairs from the positive training data, and the – label for the pairs sample from the negative data.

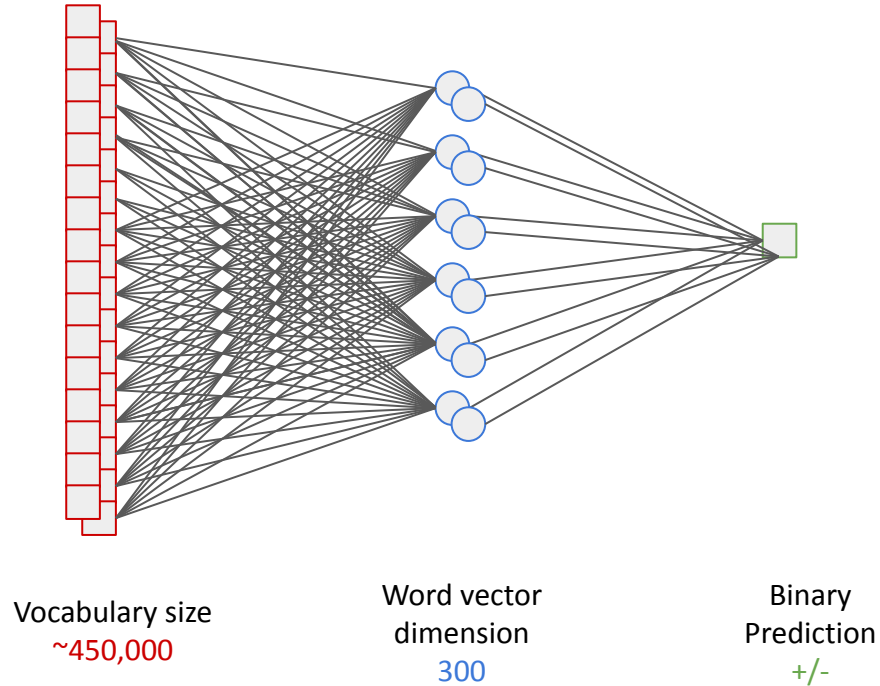
Focusing on one target word t

$$\begin{aligned} L(\theta) &= \log P(+|t, c) + \sum_{i=1}^k \log P(-|t, n_i) \\ &= \log \sigma(c \cdot t) + \sum_{i=1}^k \log \sigma(-n_i \cdot t) \\ &= \log \frac{1}{1+e^{-c \cdot t}} + \sum_{i=1}^k \log \frac{1}{1+e^{n_i \cdot t}} \end{aligned}$$

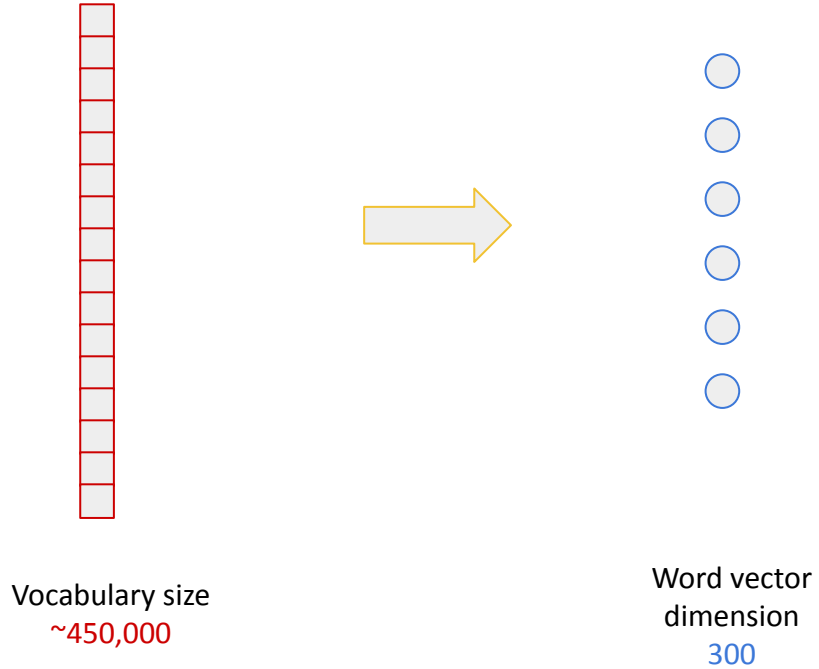
Training step (via GD)



Word2vec – Forward Step



Word2Vec – Result



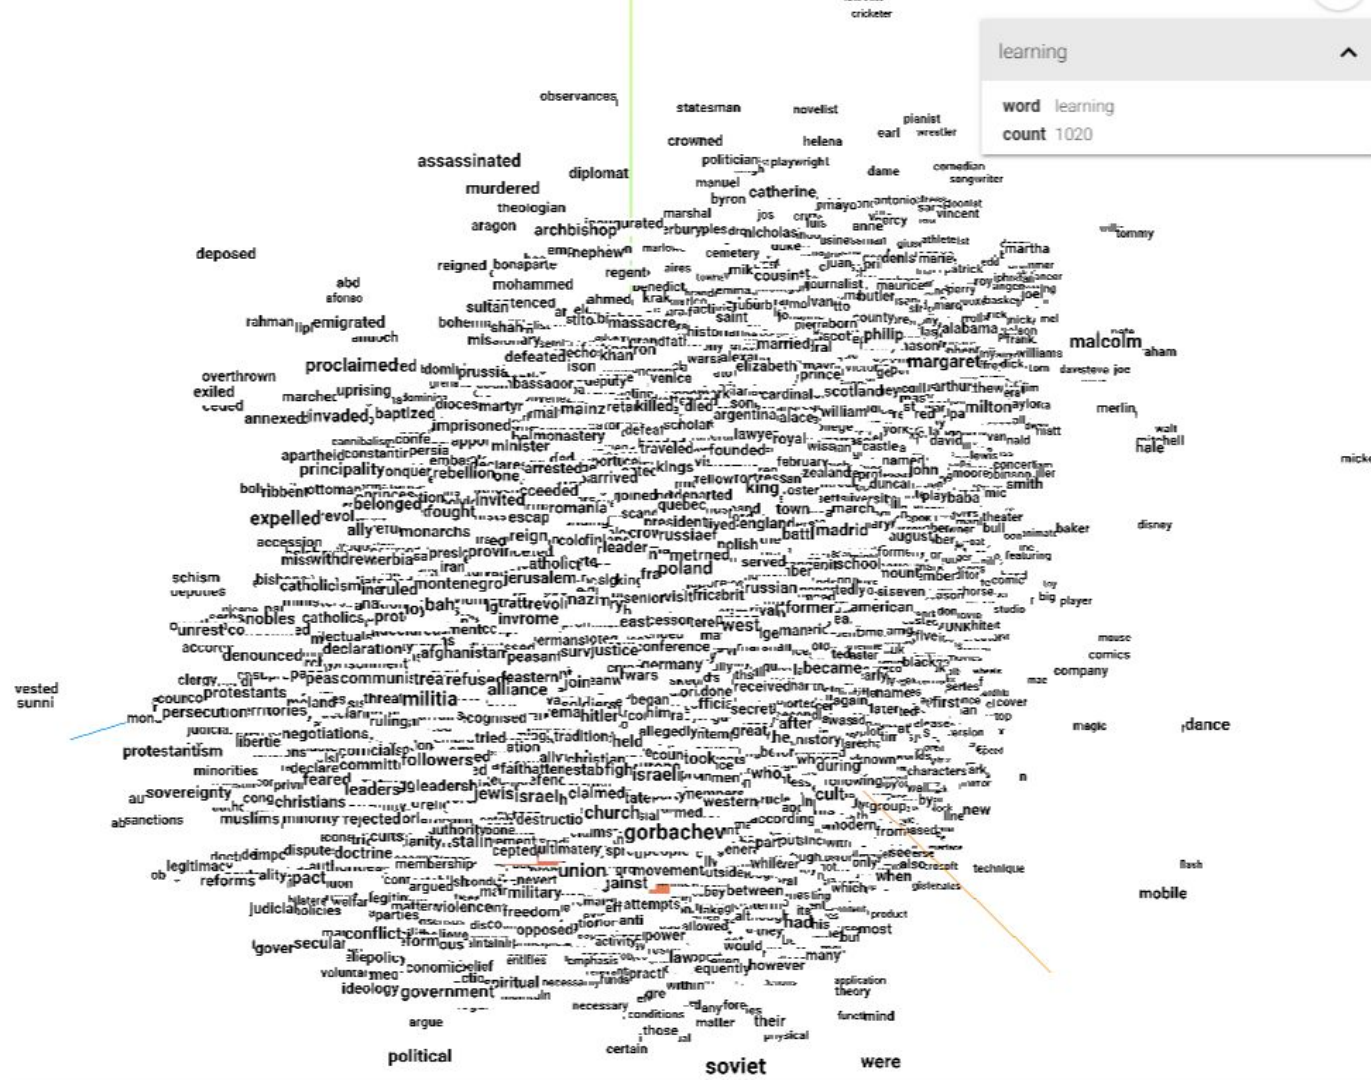
Summary: How to learn word2vec (skip-gram) embeddings

- Start with 2 random 300-dimensional layers (1 NN)
- In a corpus, take pairs of words that co-occur as positive examples
- Take pairs of words that don't as negative examples
- Train a classifier to distinguish these by slowly adjusting all the embeddings to improve the classifier performance
- Throw away the classifier code and keep the embeddings.

Evaluating embeddings

Compare to human scores on word similarity-type tasks:

- WordSim-353 (Finkelstein et al., 2002)
- **SimLex-999 (Hill et al., 2015)**
- Stanford Contextual Word Similarity (SCWS) dataset (Huang et al., 2012)
- TOEFL dataset: *Levied is closest in meaning to: imposed, believed, requested, correlated*



Search by word

neighbors 20

distance COSINE EUCLIDEAN

Nearest points in the original space:

| | |
|---------------|-------|
| teaching | 0.299 |
| educational | 0.356 |
| knowledge | 0.365 |
| learn | 0.383 |
| communication | 0.383 |
| cognitive | 0.386 |
| education | 0.387 |
| understanding | 0.409 |
| study | 0.415 |
| learned | 0.417 |
| mathematics | 0.422 |
| thinking | 0.426 |
| studying | 0.431 |
| skills | 0.433 |
| computational | 0.434 |
| logic | 0.436 |
| teachers | 0.437 |
| experience | 0.443 |
| tesch | 0.443 |
| memory | 0.444 |

TensorFlow Embedding Projector
<https://projector.tensorflow.org/>

Word Vector Space Models - Discussion

- The model is derived from context, but not context sensitive:
One representation (one vector) for the words: **play** and **bank**.
Yet they have many different senses, which are context dependent.
bank: border of the lake, financial institute, to pile up...
play: to play a piano, to play with a ball, a play...
- Vectors represent words, but how can we represent phrases and sentences?
- What about combining different modalities such as visual and textual ?

tiger



rose



zebra



Today's Agenda

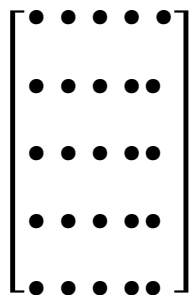
- Word2Vec
- **GloVe**
- Sentence representations
- Applications

GloVe

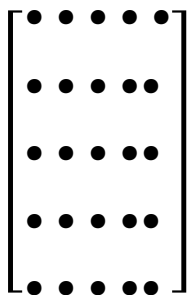
GloVe - TL;DR

- GloVe is a static word embedding algorithm
- Combines ideas from co-occurrence and language modeling
- Essentially closed the static word embedding lit.
 - Whenever we say we use *static word embedding*, we use GloVe

Word2vec parameters and computations



U
outside



V
center



$U \cdot v_4^T$
dot product



$\text{softmax}(U \cdot v_4^T)$
probabilities

"Bag of words" model!

→ The model makes the same predictions at each position

We want a model that gives a reasonably high probability estimate to *all* words that occur in the context (at all often)

Gradient Descent

- Update equation (in matrix notation):

$$\theta^{new} = \theta^{old} - \alpha \nabla_{\theta} J(\theta)$$

α = *step size* or *learning rate*

- Update equation (for a single parameter):

$$\theta_j^{new} = \theta_j^{old} - \alpha \frac{\partial}{\partial \theta_j^{old}} J(\theta)$$

- Algorithm:

```
while True:
    theta_grad = evaluate_gradient(J, corpus, theta)
    theta = theta - alpha * theta_grad
```

Stochastic Gradient Descent

- **Problem:** $J(\theta)$ is a function of **all** windows in the corpus (often, billions!)
 - So $\nabla_{\theta} J(\theta)$ is **very expensive to compute**
- You would wait a very long time before making a single update!
- **Very** bad idea for pretty much all neural nets!
- **Solution: Stochastic gradient descent (SGD)**
 - Repeatedly sample windows, and update after each one, or each small batch
- Algorithm:

```
while True:
    window = sample_window(corpus)
    theta_grad = evaluate_gradient(J, window, theta)
    theta = theta - alpha * theta_grad
```

Stochastic gradients with word vectors! [Aside]

- Iteratively take gradients at each such window for SGD
- But in each window, we only have at most $2m + 1$ words, so $\nabla_{\theta} J_t(\theta)$ is very sparse!

$$\nabla_{\theta} J_t(\theta) = \begin{bmatrix} 0 \\ \vdots \\ \nabla_{v_{like}} \\ \vdots \\ 0 \\ \nabla_{u_I} \\ \vdots \\ \nabla_{u_{learning}} \\ \vdots \end{bmatrix} \in \mathbb{R}^{2dV}$$

Stochastic gradients with word vectors!

- We might only update the word vectors that actually appear!
- Solution: either you need sparse matrix update operations to only update certain **rows** of full embedding matrices U and V , or you need to keep around a hash for word vectors

$$|V| \begin{bmatrix} \bullet & \bullet & \bullet & \bullet & \bullet \\ \bullet & \bullet & \bullet & \bullet & \bullet \\ \bullet & \bullet & \bullet & \bullet & \bullet \\ \bullet & \bullet & \bullet & \bullet & \bullet \end{bmatrix}^d$$

- If you have millions of word vectors and do distributed computing, it is important to not have to send gigantic updates around!

Why not capture co-occurrence counts directly?

Building a co-occurrence matrix X

- 2 options: windows vs. full document
- Window: Similar to word2vec, use window around each word
 - captures some syntactic and semantic information
- Word-document co-occurrence matrix will give general topics (all sports terms will have similar entries) leading to “Latent Semantic Analysis”

Example: Window based co-occurrence matrix

- Window length 1 (more common: 5–10)
- Symmetric (irrelevant whether left or right context)
- Example corpus:
 - I like deep learning
 - I like NLP
 - I enjoy flying

| counts | I | like | enjoy | deep | learning | NLP | flying | . |
|----------|---|------|-------|------|----------|-----|--------|---|
| I | 0 | 2 | 1 | 0 | 0 | 0 | 0 | 0 |
| like | 2 | 0 | 0 | 1 | 0 | 1 | 0 | 0 |
| enjoy | 1 | 0 | 0 | 0 | 0 | 0 | 1 | 0 |
| deep | 0 | 1 | 0 | 0 | 1 | 0 | 0 | 0 |
| learning | 0 | 0 | 0 | 1 | 0 | 0 | 0 | 1 |
| NLP | 0 | 1 | 0 | 0 | 0 | 0 | 0 | 1 |
| flying | 0 | 0 | 1 | 0 | 0 | 0 | 0 | 1 |
| . | 0 | 0 | 0 | 0 | 1 | 1 | 1 | 0 |

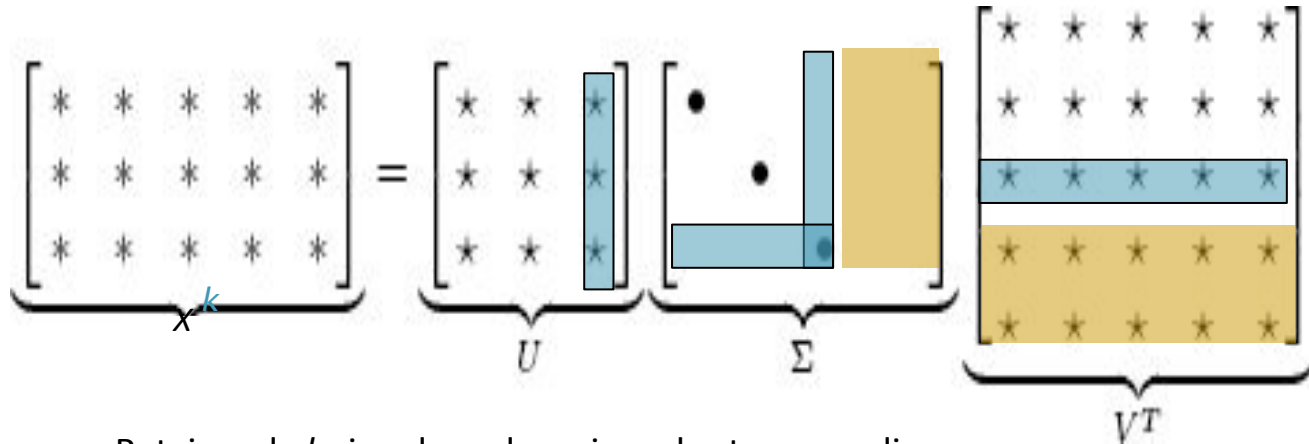
Co-occurrence vectors

- Simple count co-occurrence vectors
 - Vectors increase in size with vocabulary
 - Very high dimensional: require a lot of storage (though sparse)
 - Subsequent classification models have sparsity issues -> Models are less robust
- Low-dimensional vectors
 - Idea: store “most” of the important information in a fixed, small number of dimensions: a dense vector
 - Usually 25–1000 dimensions, similar to word2vec
 - How to reduce the dimensionality?

Classic Method: Dimensionality Reduction on X

Singular Value Decomposition of co-occurrence matrix X

Factorizes X into $U\Sigma V^T$, where U and V are orthonormal



Retain only k singular values, in order to generalize.

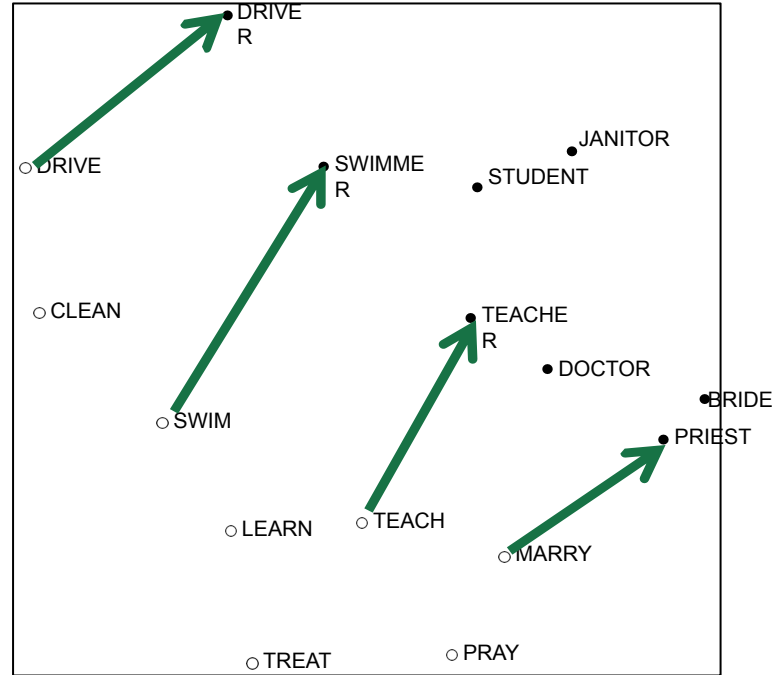
X is the best rank k approximation to X , in terms of least squares.

Classic linear algebra result. Expensive to compute for large matrices.

Hacks to X (several used in Rohde et al. 2005 in COALS)

- Running an SVD on raw counts doesn't work well
- Scaling the counts in the cells can help ***a lot***
 - Problem: function words (*the, he, has*) are too frequent → syntax has too much impact. Some fixes:
 - log the frequencies
 - $\min(X, t)$, with $t \approx 100$
 - Ignore the function words
- Ramped windows that count closer words more than further away words
- Use Pearson correlations instead of counts, then set negative values to 0
- Etc.

Interesting semantic patterns emerge in the scaled vectors



COALS model from
Rohde et al. ms., 2005. An Improved Model of Semantic Similarity Based on Lexical Co-Occurrence

Towards GloVe: Count based vs. direct prediction

- LSA, HAL
- COALS, Hellinger-PCA

- Fast training
- Efficient usage of statistics
- Primarily used to capture word similarity
- Disproportionate importance given to large counts

- Skip-gram/CBOW
- NNLM, HLBL, RNN

- Scales with corpus size
- Inefficient usage of statistics
- Generate improved performance on other tasks
- Can capture complex patterns beyond word similarity

Encoding meaning components in vector differences

Crucial insight: Ratios of co-occurrence probabilities can encode meaning components

| | $x = \text{solid}$ | $x = \text{gas}$ | $x = \text{water}$ | $x = \text{random}$ |
|---|--------------------|------------------|--------------------|---------------------|
| $P(x \text{ice})$ | large | small | large | small |
| $P(x \text{steam})$ | small | large | large | small |
| $\frac{P(x \text{ice})}{P(x \text{steam})}$ | large | small | ~ 1 | ~ 1 |

Encoding meaning in vector differences

Crucial insight: Ratios of co-occurrence probabilities can encode meaning components

| | $x = \text{solid}$ | $x = \text{gas}$ | $x = \text{water}$ | $x = \text{fashion}$ |
|---|--|--|--|----------------------|
| $P(x \text{ice})$ | 1.9×10^{-4} | 6.6×10^{-5} | 3.0×10^{-3} | 1.7×10^{-5} |
| $P(x \text{steam})$ | 2.2×10^{-5} | 7.8×10^{-4} | 2.2×10^{-3} | 1.8×10^{-5} |
| $\frac{P(x \text{ice})}{P(x \text{steam})}$ | 8.9 | 8.5×10^{-2} | 1.36 | 0.96 |

Encoding meaning in vector differences

Q: How can we capture ratios of co-occurrence probabilities as linear meaning components in a word vector space?

A: Log-bilinear model:
with vector differences

$$w_i \cdot w_j = \log P(i|j)$$

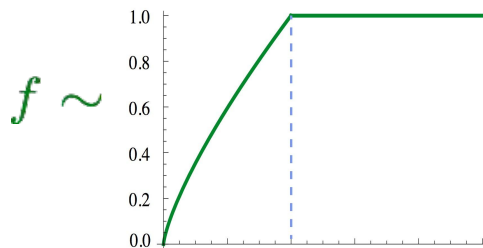
$$w_x \cdot (w_a - w_b) = \log \frac{P(x|a)}{P(x|b)}$$

Combining the best of both worlds - GloVe

$$w_i \cdot w_j = \log P(i|j)$$

$$J = \sum_{i,j=1}^V f(X_{ij}) \left(w_i^T \tilde{w}_j + b_i + \tilde{b}_j - \log X_{ij} \right)^2$$

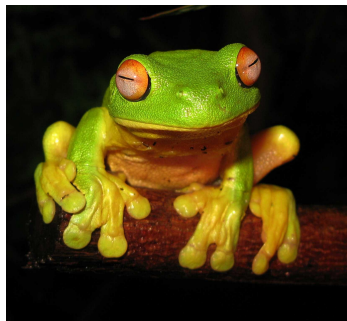
- Fast training
- Scalable to huge corpora
- Good performance even with small corpus and small vectors



GloVe results

Nearest words to **frog**:

1. frogs
2. toad
3. litoria
4. leptodactylidae
5. rana
6. lizard
7. eleutherodactylus



litoria



leptodactylidae



rana



eleutherodactylus

How to evaluate word vectors?

- Related to general evaluation in NLP: Intrinsic vs. extrinsic
- Intrinsic:
 - Evaluation on a specific/intermediate subtask
 - Fast to compute
 - Helps to understand that system
 - Not clear if really helpful unless correlation to real task is established
- Extrinsic:
 - Evaluation on a real task
 - Can take a long time to compute accuracy
 - Unclear if the subsystem is the problem or its interaction or other subsystems
 - If replacing exactly one subsystem with another improves accuracy

Intrinsic word vector evaluation

- Word Vector Analogies

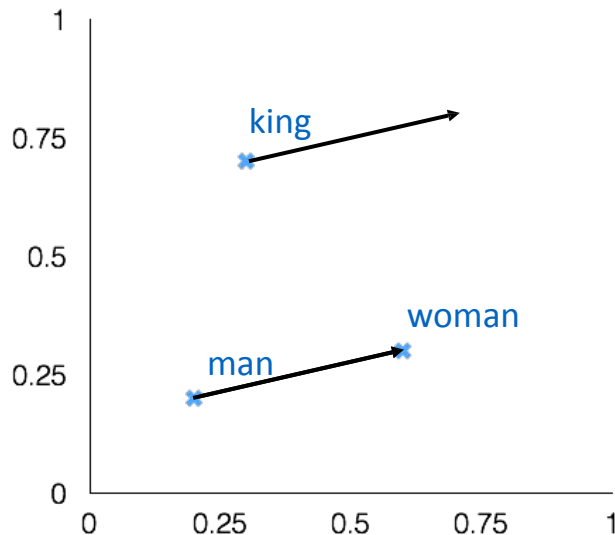
a:b :: c:?

man:woman :: king:?

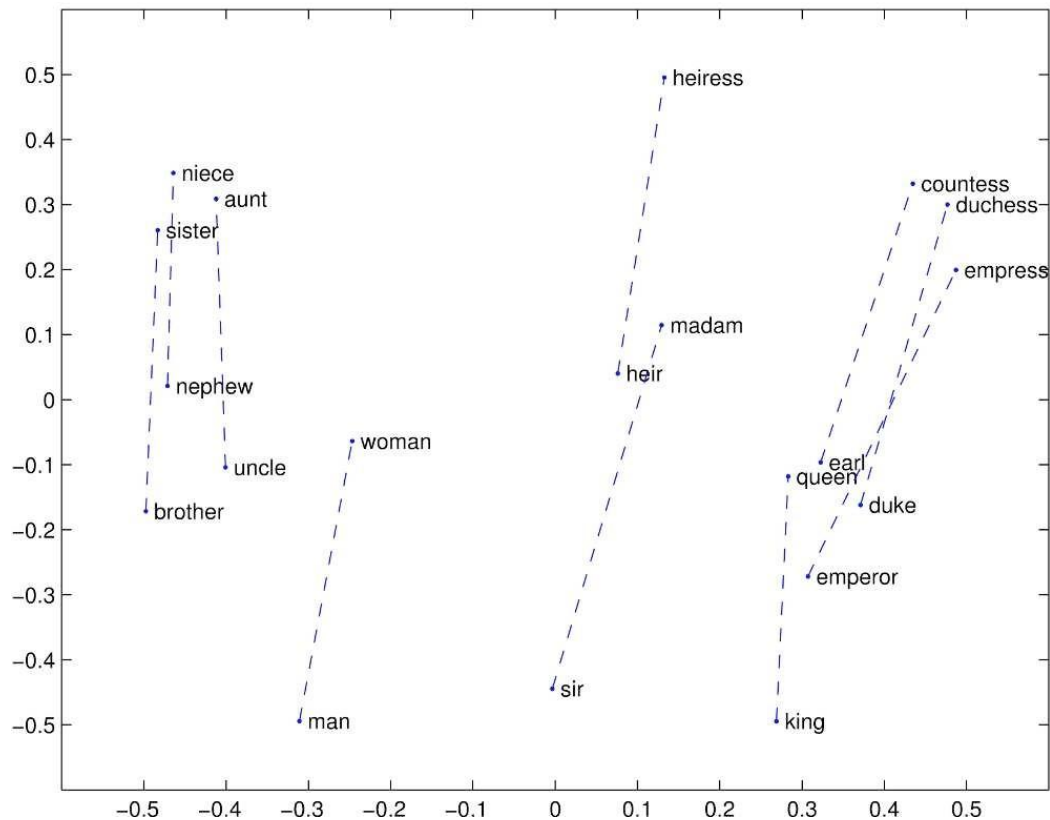


$$d = \arg \max_i \frac{(x_b - x_a + x_c)^T x_i}{\|x_b - x_a + x_c\|}$$

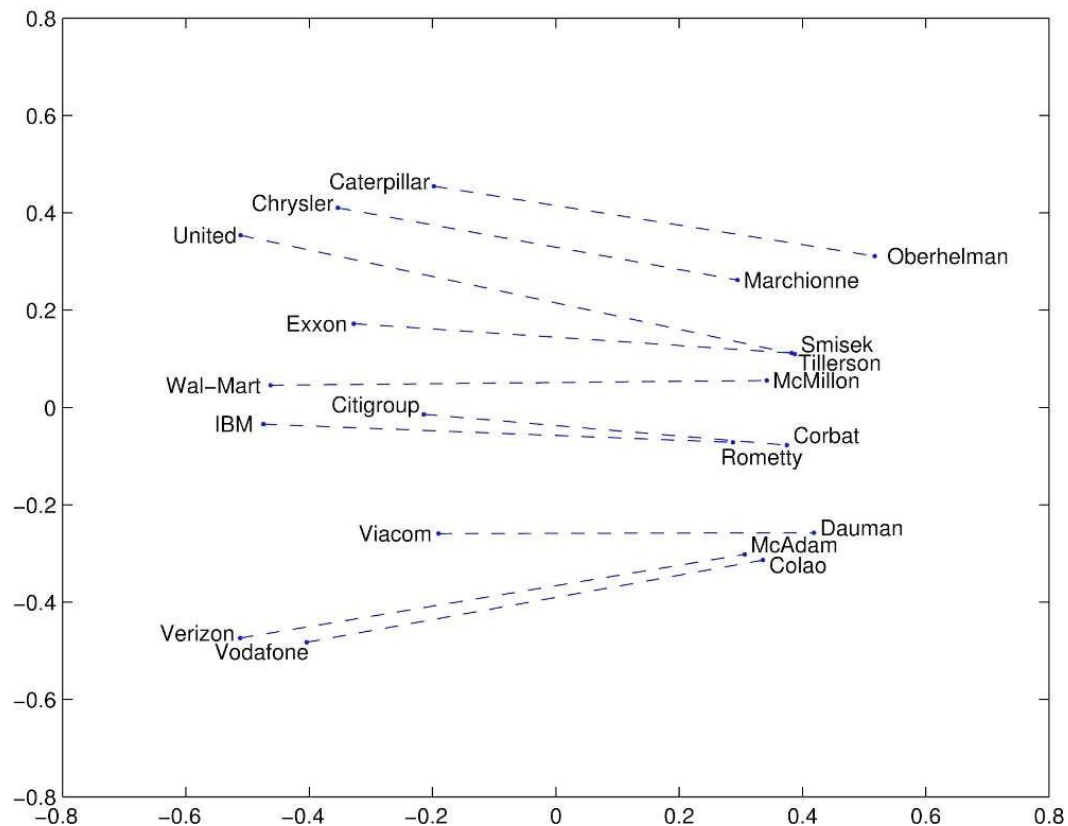
- Evaluate word vectors by how well their cosine distance after addition captures intuitive semantic and syntactic analogy questions
- Discarding the input words from the search!
- Problem: What if the information is there but not linear?



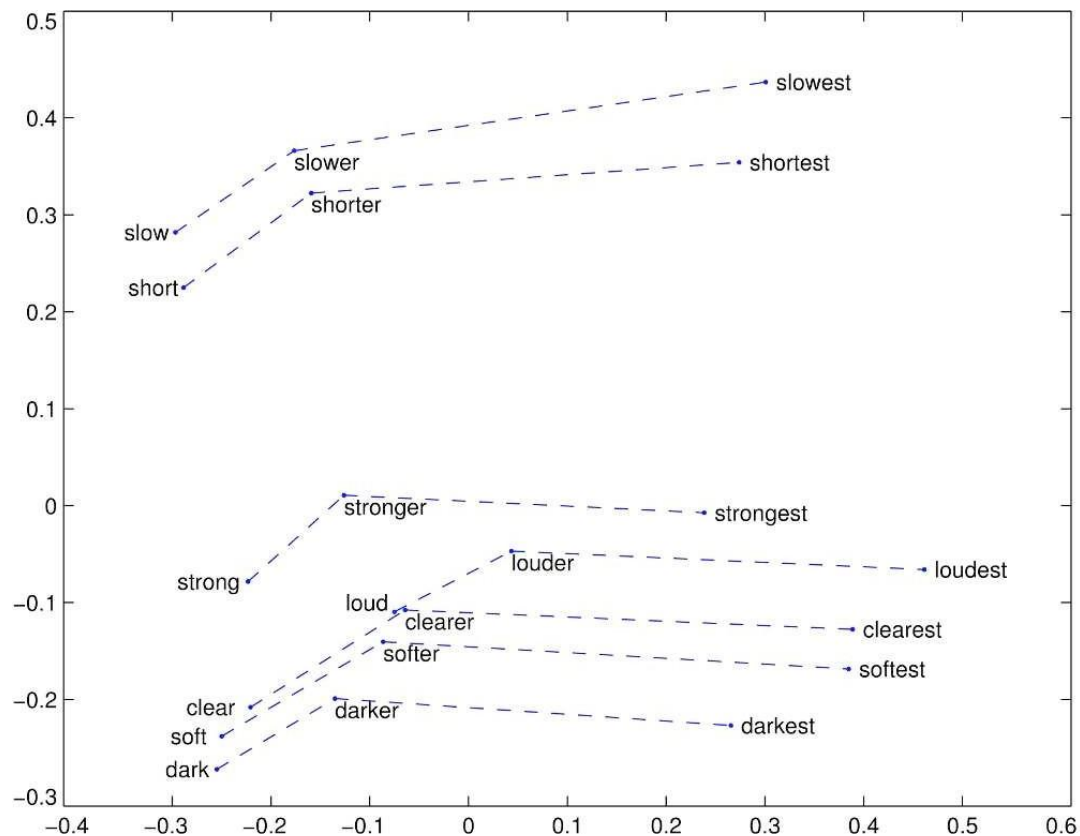
Glove Visualizations



Glove Visualizations: Company - CEO



Glove Visualizations: Comparatives and Superlatives



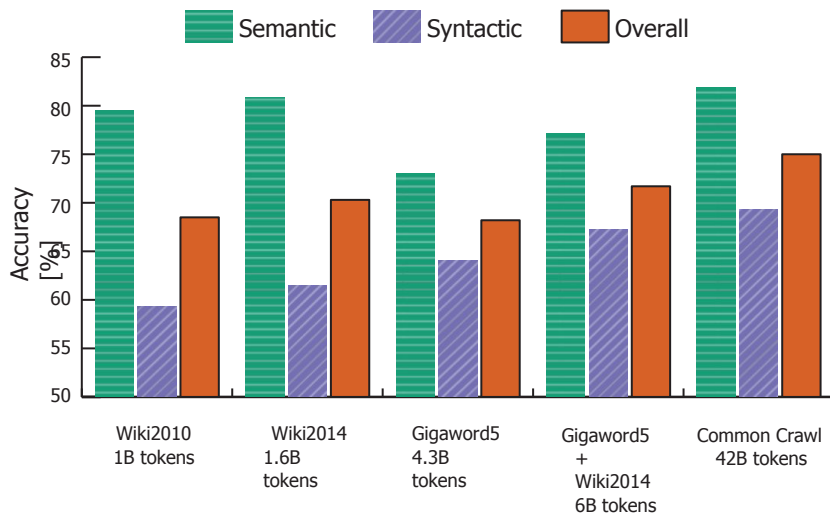
Analogy evaluation and hyperparameters

Glove word vectors evaluation

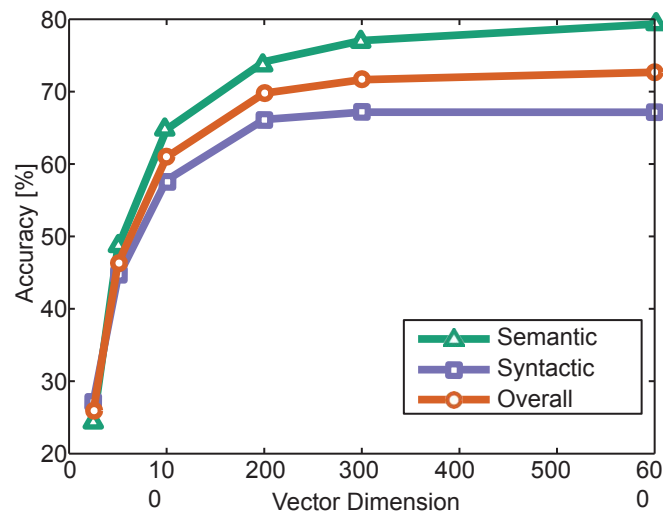
| Model | Dim. | Size | Sem. | Syn. | Tot. |
|-------------------|------|------|-------------|-------------|-------------|
| SVD | 300 | 6B | 6.3 | 8.1 | 7.3 |
| SVD-S | 300 | 6B | 36.7 | 46.6 | 42.1 |
| SVD-L | 300 | 6B | 56.6 | 63.0 | 60.1 |
| CBOW [†] | 300 | 6B | 63.6 | <u>67.4</u> | 65.7 |
| SG [†] | 300 | 6B | 73.0 | 66.0 | 69.1 |
| GloVe | 300 | 6B | <u>77.4</u> | 67.0 | <u>71.7</u> |

Analogy evaluation and hyperparameters

- More data helps
- Wikipedia is better than news text!



- Dimensionality
- Good dimension is ~300



Another intrinsic word vector evaluation

- Word vector distances and their correlation with human judgments
- Example dataset: SimLex999 <http://www.cl.cam.ac.uk/~fh295/simlex.html>

| Word 1 | Word 2 | Human (mean) |
|-----------|----------|--------------|
| tiger | cat | 7.35 |
| tiger | tiger | 10 |
| book | paper | 7.46 |
| computer | internet | 7.58 |
| plane | car | 5.77 |
| professor | doctor | 6.62 |
| stock | phone | 1.62 |
| stock | CD | 1.31 |
| stock | jaguar | 0.92 |

Correlation evaluation

- Word vector distances and their correlation with human judgments

| Model | Size | WS353 | MC | RG | SCWS | RW |
|-------------------|------|-------------|-------------|-------------|-------------|-------------|
| SVD | 6B | 35.3 | 35.1 | 42.5 | 38.3 | 25.6 |
| SVD-S | 6B | 56.5 | 71.5 | 71.0 | 53.6 | 34.7 |
| SVD-L | 6B | 65.7 | <u>72.7</u> | 75.1 | 56.5 | 37.0 |
| CBOW [†] | 6B | 57.2 | 65.6 | 68.2 | 57.0 | 32.5 |
| SG [†] | 6B | 62.8 | 65.2 | 69.7 | <u>58.1</u> | 37.2 |
| GloVe | 6B | <u>65.8</u> | <u>72.7</u> | <u>77.8</u> | 53.9 | <u>38.1</u> |
| SVD-L | 42B | 74.0 | 76.4 | 74.1 | 58.3 | 39.9 |
| GloVe | 42B | <u>75.9</u> | <u>83.6</u> | <u>82.9</u> | <u>59.6</u> | <u>47.8</u> |
| CBOW [*] | 100B | 68.4 | 79.6 | 75.4 | 59.4 | 45.5 |

Extrinsic word vector evaluation

- Extrinsic evaluation of word vectors: All subsequent NLP tasks in this class. More examples soon.
- One example where good word vectors should help directly: **named entity recognition**: identifying references to a person, organization or location

| Model | Dev | Test | ACE | MUC7 |
|----------|------|------|------|------|
| Discrete | 91.0 | 85.4 | 77.4 | 73.4 |
| SVD | 90.8 | 85.7 | 77.3 | 73.7 |
| SVD-S | 91.0 | 85.5 | 77.6 | 74.3 |
| SVD-L | 90.5 | 84.8 | 73.6 | 71.5 |
| HPCA | 92.6 | 88.7 | 81.7 | 80.7 |
| HSMN | 90.5 | 85.7 | 78.7 | 74.7 |
| CW | 92.2 | 87.4 | 81.7 | 80.2 |
| CBOW | 93.1 | 88.2 | 82.2 | 81.1 |
| GloVe | 93.2 | 88.3 | 82.9 | 82.2 |

Today's Agenda

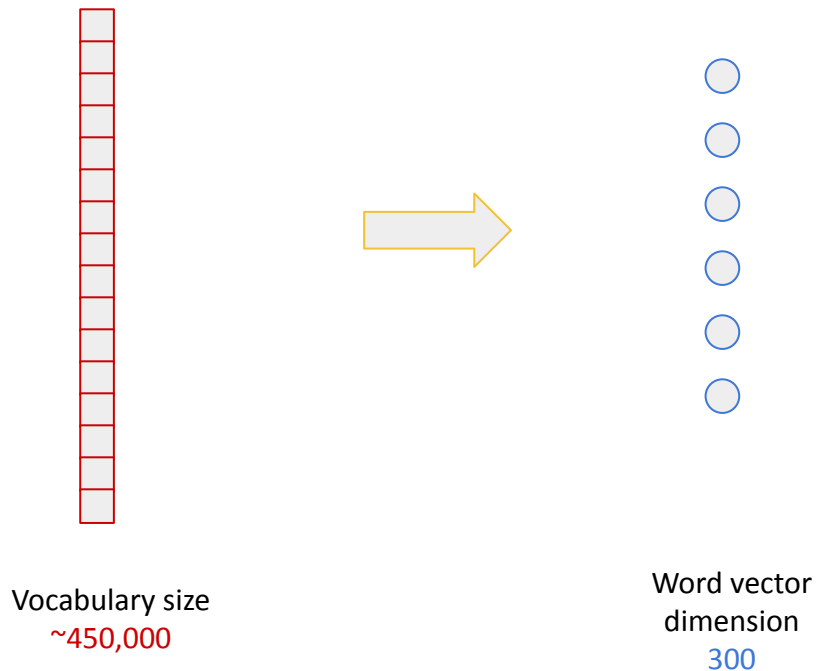
- Word2Vec
- GloVe
- **Sentence representations**
- Applications

Sentence representations

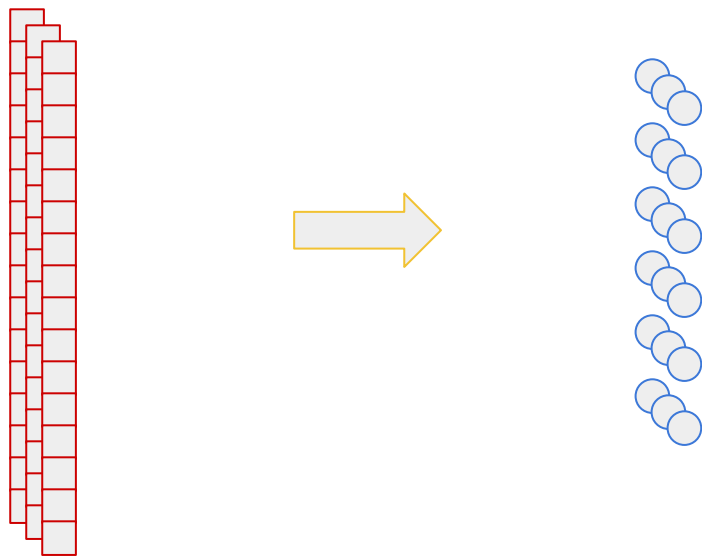
A sentence is a bunch of words

- We learned how to represent words
- We can just concatenate
- But what will we lose?

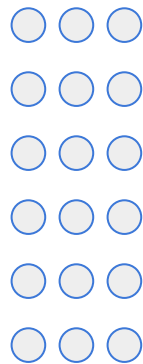
Words now represent meaning



Sentence = Concatenated Word Rep.



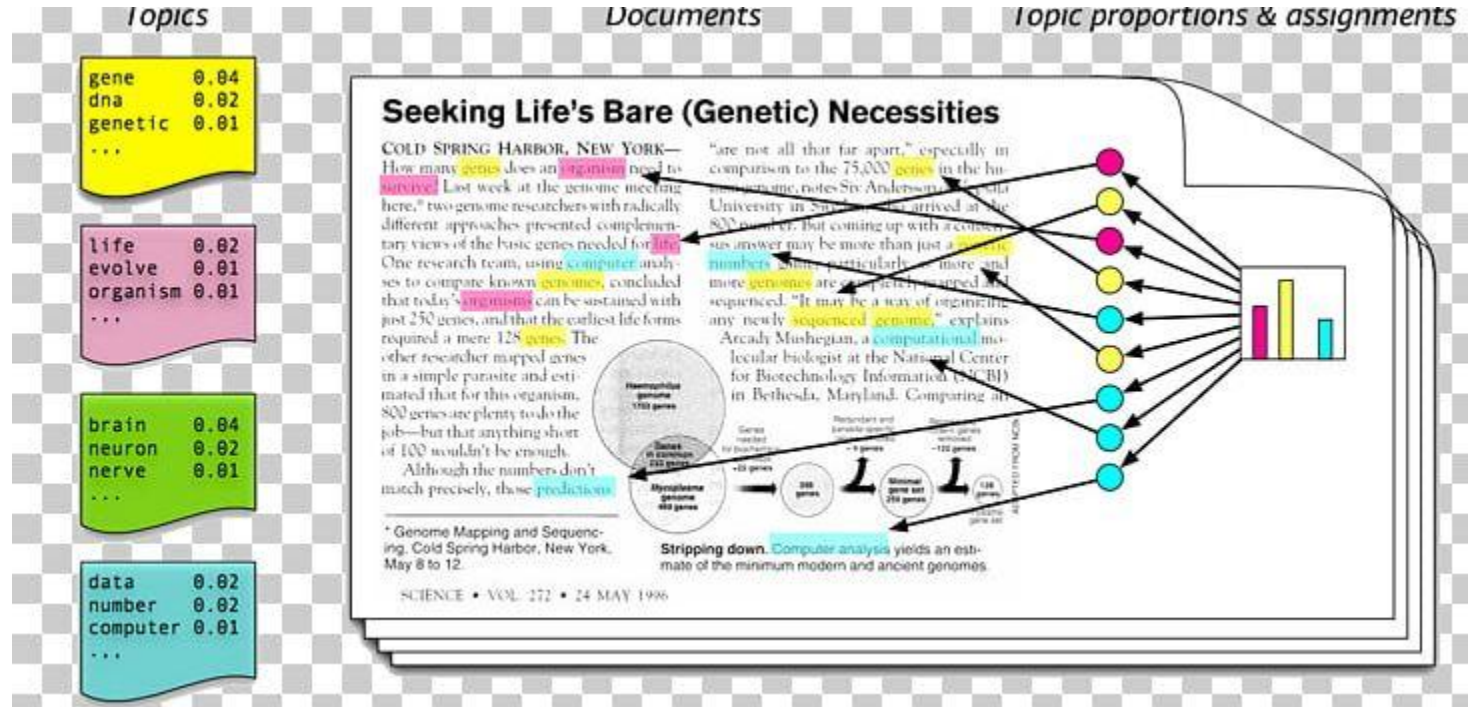
Concatenated Word Rep.



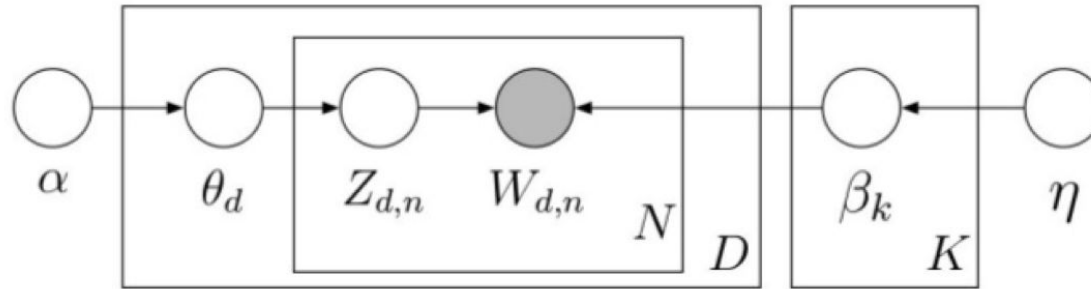
Count-based Representation

- We also learned to represent windows with multiple words
- But this is high-dimensional and sparse
- We can use only some of the dimensions
- We can cluster, but how?

Latent Dirichlet Allocation



Latent Dirichlet Allocation



K – total number of topics

β_k – topic, a distribution over the vocabulary

D – total number of documents

θ_d – per-document topic proportions

N – total number of words in a document (it fact, it should be N_d)

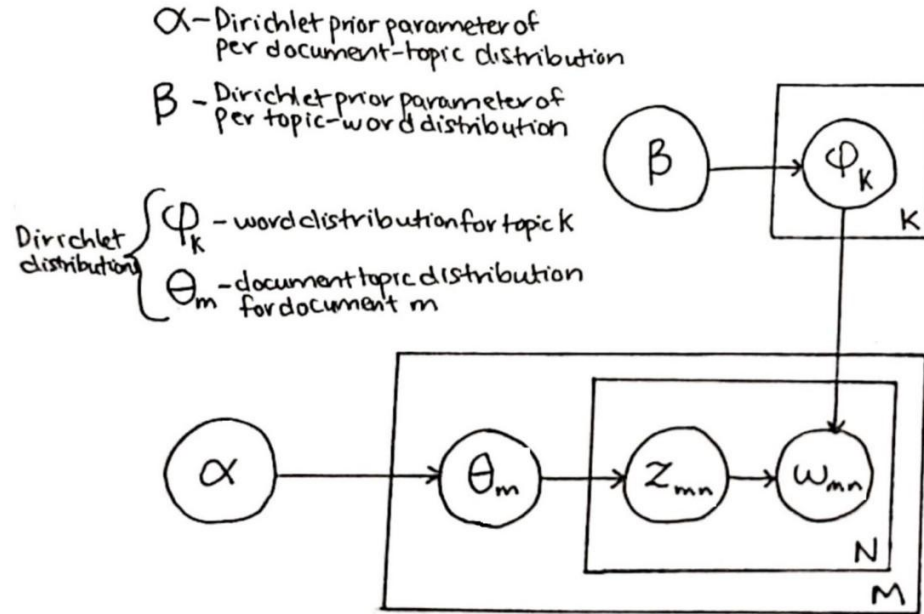
$Z_{d,n}$ – per-word topic assignment

$W_{d,n}$ – observed word

α, η – Dirichlet parameters

- Several **inference algorithms** are available (e.g. sampling based)
- A few **extensions** to LDA were created:
 - Bigram Topic Model

Latent Dirichlet Allocation



K - number of topics
 N - number of words in document
 M - number of documents

z_{mn} - word-topic assignment for w_{mn}
 w_{mn} - observed word i.e. the n -th word in the m -th document. } multinomial distributions

Today's Agenda

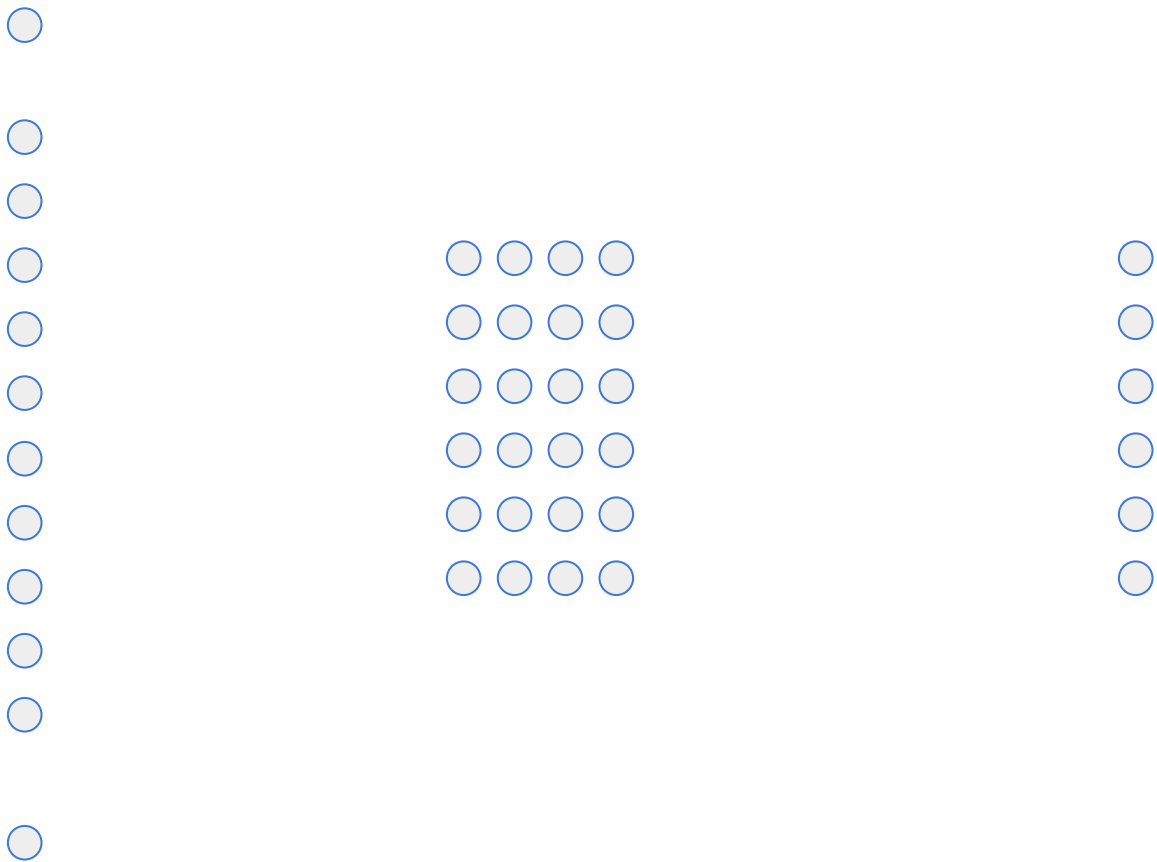
- Word2Vec
- GloVe
- Sentence representations
- **Applications**

Applications

Let's build some models together (on board)

- Sentiment Classifier
- Named Entity Recognizer
- Part-of-Speech Tagger

Sentence Representation



Sentiment Classification

| | text | sentiment |
|---|---|------------------|
| 0 | For a movie that gets no respect there sure ar... | 0 |
| 1 | Bizarre horror movie filled with famous faces ... | 0 |
| 2 | A solid, if unremarkable film. Matthau, as Ein... | 0 |
| 3 | It's a strange feeling to sit alone in a theat... | 0 |
| 4 | You probably all already know this by now, but... | 0 |
| 5 | I saw the movie with two grown children. Altho... | 0 |
| 6 | You're using the IMDb. You've given some heft... | 0 |
| 7 | This was a good film with a powerful message o... | 0 |
| 8 | Made after QUARTET was, TRIO continued the qua... | 0 |
| 9 | For a mature man, to admit that he shed a tear... | 0 |

Named Entity Recognition

Person p Loc l Org o Event e Date d Other z

Barack Hussein Obama II * (born August 4, 1961 *) is an American * attorney and politician who served as the 44th President of the United States * from January 20, 2009 *, to January 20, 2017 *. A member of the Democratic Party *, he was the first African American * to serve as president. He was previously a United States Senator * from Illinois * and a member of the Illinois State Senate *.

Parts of Speech

8 Parts of Speech

NOUN

A **noun** names a person, place, things or idea.

Examples

dog, cat, horse, student, teacher, apple, Mary and etc...

ADVERB

An **adverb** tells how often, how, when, where. It can describe a verb, an adjective or an adverb.

Examples

loudly, always, never, late, soon etc...

VERB

A **verb** is a word or group of words that describes an action, experience.

Examples

realize, walk, see, look, sing, sit, listen and etc...

ADJECTIVE

An **adjective** describes a noun or pronoun.

Examples

red, tall, fat, long, short, blue, beautiful, sour and etc...

PREPOSITION

A **preposition** is used before a noun, pronoun, or gerund to show place, time, direction in a sentence.

Examples

at, in, on, about, to, for, from and etc...

CONJUNCTION

Conjunctions join words or groups of words in a sentence.

Examples

and, because, yet, therefore, moreover, since, or, so, until, but and etc...

PRONOUN

Pronouns replace the name of a person, place, thing or idea in a sentence.

Examples

he, she, it, we, they, him, her, this ,that and etc...

INTERJECTION

Interjections express strong emotion and is often followed by an exclamation point.

Examples

Bravo! Well! Aha! Hooray! Yeah! Oops! Phew!

REAL ENGLISH.IK